

# Prove platform judgment through the work.

An affirmative response to the role's low-code and AI delivery expectations

The relevant question is not whether a candidate can recite a tool list. It is whether the candidate can choose, configure, bound, evaluate, and operationalize the right mechanism for the workflow.

## THE HIRING LOGIC

The posting expects advanced proficiency with Microsoft Power Platform and emerging automation tools such as n8n, Make, and Zapier. Russell's verified evidence establishes deep adjacent strength—agentic workflow design, RAG, human-in-the-loop systems, APIs/databases/coding foundations, operating transformation, quality systems, stakeholder translation, and adoption—but does not substantiate production tenure in every named platform.

The right response is neither to hide that boundary nor turn it into a recruiter-facing warning label. It is to make technical transferability observable early.

## INVARIANT ARCHITECTURE RUSSELL ALREADY WORKS IN

|                                                                                                                 |                                                                                                                    |
|-----------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| <b>Workflow decomposition</b><br>Actors, inputs, decisions, actions, exceptions, consequences, and handoffs.    | <b>Context and data</b><br>Sources, retrieval, structure, freshness, grounding, privacy, and integration boundary. |
| <b>Agent behavior</b><br>Tools, planning, action permissions, human authority, escalation, fallback, and trace. | <b>Proof and operation</b><br>Evals, testing, cost/latency, adoption, support, ownership, and improvement loop.    |

## WHERE CROWE-SPECIFIC LEARNING MUST OCCUR

| Area                                            | Learning target                                                                                                               | Proof                                                    |
|-------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------|
| Power Platform architecture                     | Environment strategy, Dataverse/data model, connectors, Power Apps/Automate, AI Builder/Copilot patterns, governance and ALM. | Bounded workflow build reviewed against Crowe standards. |
| Adjacent orchestration tools                    | n8n/Make/Zapier fit, self-hosted vs. managed implications, credentials, connectors, observability, retries, maintainability.  | Platform comparison and configured integration slice.    |
| Azure AI services                               | Approved service patterns, security, model endpoints, content safety, monitoring, cost and deployment conventions.            | Architecture explanation plus test/evidence record.      |
| Consulting delivery technical value case method | Crowe estimation, documentation, design review, quality gates, client handoff and practice asset harvesting.                  | Complete workstream contribution with positive review.   |

# A 60–90 minute way to test the transfer.

## PROMPT

Choose one realistic workflow: for example, a governed operations copilot that receives an exception, retrieves the relevant policy and transaction context, recommends an action, requests human approval when required, updates a system of record, and preserves a decision trace.

## SESSION STRUCTURE

| Time      | Activity                                                                                                                              | Evidence produced                                          |
|-----------|---------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------|
| 0–15 min  | Clarify outcome, users, current work, exceptions, source systems, risk and ownership.                                                 | Workflow frame and non-goals.                              |
| 15–35 min | Compare candidate platforms against connector fit, data model, autonomy ceiling, governance, reuse, supportability and time-to-value. | Decision matrix with assumptions and unresolved questions. |
| 35–55 min | Sketch context flow, tools/actions, prompts/instructions, human checkpoints, errors, logs and fallback.                               | Solution architecture and action envelope.                 |
| 55–75 min | Define evaluation cases, acceptance criteria, adoption evidence and operating owner.                                                  | Proof plan and definition of done.                         |
| 75–90 min | Select the smallest bounded build; identify where Crowe review or pairing is needed.                                                  | Build plan, risk register and next decision.               |

## PLATFORM DECISION CRITERIA

| Integration fit                                                    | Action envelope                                                | Governance                                                   | Supportability                                                    | Value speed                                                      |
|--------------------------------------------------------------------|----------------------------------------------------------------|--------------------------------------------------------------|-------------------------------------------------------------------|------------------------------------------------------------------|
| Required connectors, APIs, events, authentication, data residency. | What the agent/ workflow can do, with what approval and trace. | Environments, access, audit, DLP, lifecycle, change control. | Observability, retry, testing, ownership, skills and vendor risk. | Time to useful slice, reuse, operating cost and adoption burden. |

## WHAT A STRONG PERFORMANCE WOULD SHOW

- Russell asks business, workflow, integration, risk, adoption, and ownership questions before selecting a tool.
- He can reason across technical and nontechnical audiences without hiding tradeoffs behind jargon.
- He distinguishes assistive, approval-gated, and autonomous action; defines human authority and failure handling.
- He proposes testable evidence and a smallest useful slice rather than a grand architecture.
- He identifies where his platform knowledge needs review and converts that need into a focused learning action.

### Hiring decision signal

The session should answer whether Russell’s operating and AI architecture depth transfers quickly into Crowe’s low-code stack—and whether he demonstrates the curiosity, ownership, judgment, and communication expected of a senior forward-deployed consultant.